

CS174A : Introduction to Computer Graphics

Royce 190
TT 4-6pm


Scott Friedman, Ph.D
UCLA Institute for Digital Research and Education

Term Project

- Details will be introduced next Tuesday.
- Things to start considering now are:
 - Teams will be a minimum of *three* people.
 - Use the class forum to find partners if you need to.
 - Teams can have up to *five* people in your group.
 - Your team will submit a detailed project proposal.
 - “we are going to write a game”, will not do for a proposal.
 - The proposals will be due the following week.
 - The TAs and I will review them and make comments.

Trackball


- Last time we talked about rotation...
 - A useful/intuitive interaction technique.
 - Plant a unit hemisphere into the x - z plane.
 - Using mouse positions we can reconstruct y
 - Because, $x^2 + y^2 + z^2 = 1$, $y = \sqrt{1 - x^2 - z^2}$



Trackball

- That's nice, we know $y...$
 - Now we can track, as the mouse moves, a position p_1 to p_2 on the surface of the hemisphere.
 - What we really want to do is rotate in the direction of the arc swept out from p_1 to p_2 .
 - That axis of rotation can be found via the cross product of p_1 to p_2 , the *normal*.

$$n = p_1 \times p_2$$



Trackball

- That's nice, we know n ...
 - Conveniently, n can tell us the angle between p_1 to p_2 as well because we are using a *unit* hemisphere.
$$|\sin \theta| = |n|$$
 - Now we know an angle and a vector around which the rotation is supposed to occur.
 - The book mentions a nice trick when animating the rotation in small increments – and that is to recognize that for small
 - and you can avoid the inverse sine operation.

$$\theta \approx \sin \theta$$

Trackball

- A side note...
 - When animating a rotation in small increments
 - A lot of trigonometry is involved = slow
 - Helpful to recognize that for small $\theta \approx \sin \theta$
 - and then you can avoid the inverse sine operation.
 - This implies that, *for small angles* we can use
 - Another graphics “trick”

$$R = R_z(\psi)R_y(\phi)R_x(\theta) \approx \begin{bmatrix} 1 & -\psi & \phi & 0 \\ \psi & 1 & -\theta & 0 \\ -\phi & \theta & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Trackball

- That's nice, we know the *angle*...
 - But, we just know how to rotate around x , y and z !
 - Yes, but if we align the rotation vector with, say, the z -axis we perform our rotation. *Simple!* ☺

$$R = R_x(-\theta_x)R_y(-\theta_y)R_z(\theta_z)R_y(\theta_y)R_x(\theta_x)$$


- Ugh, we don't know $R_x(\theta_x)$ or $R_y(\theta_y)$ even if we decided that $R_z(\theta_z)$ was the rotation we wanted.
- Yes, but let's understand what is going on first.

Trackball

- How can we find the x and y rotations?
 - First we need to rotate around the x -axis onto the x - z plane.
 - Recall that $\cos \theta_x = \alpha_x$
 - Then,

$$d = \sqrt{\alpha_y^2 + \alpha_z^2}$$


$$R_x(\theta_x) = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & \alpha_z / d & -\alpha_y / d & 0 \\ 0 & \alpha_y / d & \alpha_z / d & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$



Trackball

- Now we need the y rotation?
 - We can follow the same process
 - Again, recall that $\cos \theta_y = \alpha_y$
- Then,

$$R_y(\theta_y) = \begin{bmatrix} \alpha_z & 0 & -\alpha_x & 0 \\ 0 & 1 & 0 & 0 \\ \alpha_x & 0 & \alpha_z & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$



Trackball

- Actually, we can now rotate about *any* vector v fixed at a point p .

- Concatenating into M

$$M = T(p)R_x(-\theta_x)R_y(-\theta_y)R_z(\theta_z)R_y(\theta_y)R_x(\theta_x)T(-p)$$

- Now we can rotate our trackball vector!
 - Our point p is the **origin** and our vector is n .
 - This is a lot of work – is there a better way?

Quaternions

- Same result with less computation.
 - A quaternion has the form $q = w + xi + yj + zk$.
 - The terms i, j and k are imaginary.
 - Fortunately, we can ignore this fact in this class.
 - But, they are what ultimately make quaternions work.
 - Lets consider them this way $q(w, x, y, z)$
 - Lets make w the rotation
 - Lets make x, y and z be the rotation vector.

Quaternions

- Same result with less computation.
 - It is *critical* that q be *normalized*, i.e. $q=|q^2|$.
 - Or this does not work.
 - The resulting rotation matrix is

$$Q = \begin{bmatrix} 1 - 2(y^2 + z^2) & 2(xy - wz) & 2(xz + wy) & 0 \\ 2(xy + wz) & 1 - 2(x^2 + z^2) & 2(yz - wx) & 0 \\ 2(xz - wy) & 2(yz + wx) & 1 - 2(x^2 + y^2) & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Quaternions

- Same result with less computation.
 - A very nice feature of quaternions is that they allow for straightforward smooth interpolation.
 - You do this with a current rotation R and an increment R_I
 - R starts with some initial/or no rotation and rotation vector
 - R_I has a small rotation and the same rotation vector.
 - Need to *renormalize* R occasionally to keep computation stable.



$$R = R R_I$$

Viewing

- We are going to concern ourselves with two types of *viewing*.
 - Perspective viewing
 - Parallel viewing



Viewing

- In both cases we have
 - Objects,
 - Projection lines
 - Projection plane
 - Eye (center of projection)



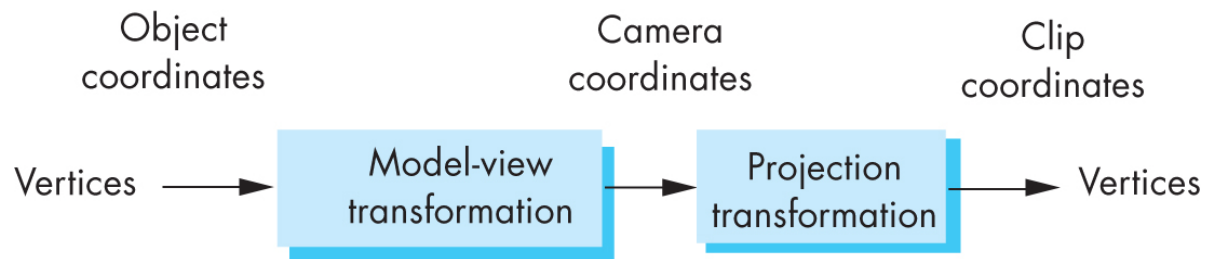
Viewing

- Our goal is to
 - Use transformations to project the vertices of objects onto the projection plane.
 - Specifically we will create transformations to go from object to camera to clip coordinates.



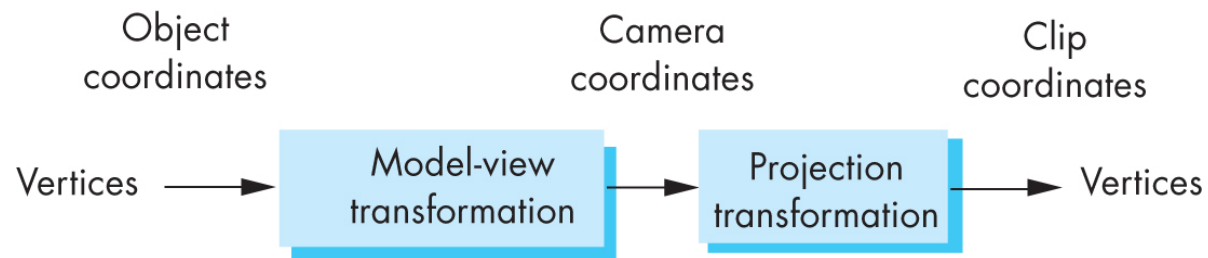
Viewing

- Previously
 - We used the default canonical view volume.
 - Which exists in clip coordinates, i.e. $(-1,1),(-1,1),(-1,1)$.
 - Last time we saw how transformations can be combined to bring objects into camera (world) coordinates
 - Collectively, the *model-view transformation*.



Viewing

- Model-view transformation
 - Does not take us all the way to clip coordinates.
 - we need a *projection transformation* for that.
 - Model-view gets objects in front of the camera, potentially.
 - A Projection defines which and how those objects will appear on the screen.



Instancing


- Useful for Assignment #2:
 - Take a single geometric object.
 - Clone it with only transformations (and possibly state).
 - Define the geometry of a cube (once).
 - Instance the cube by setting a transformation and setting some state (e.g. color) and drawing it.
 - Instance another cube by setting another transformation and setting state(e.g. color and scale) and drawing the *same geometry*.

Viewing

- Positioning the (getting things in from of the) “camera”
 - Recall that the default is “looking” down the $-z$ axis at the origin $(0,0,0)$.
 - This is equivalent to model-view set to the identity matrix.
 - Remember, transformations are specified in *reverse*.
 - That means we specify the position of the camera first.
 - We are going to look at two methods
 - Look-at
 - Yaw, pitch and roll (euler angles)

Viewing

- Look-at
 - We define three terms
 - A point describing the location of the *eye*.
 - A point the eye is looking *at*.
 - An *up* direction for the camera.



Viewing

- Look-at
 - The *at* and *eye* points give us
 - the *view-plane-normal* or *vpn*
 - the *up* vector is usually (0, 1, 0)
 - Or, (0, 1, 0, 0) in homogeneous coordinates!
 - We then calculate the following

$$\begin{aligned}vpn &= at - eye & u &= \frac{up \times n}{|up \times n|} \\ n &= \frac{vpn}{|vpn|} & v &= \frac{n \times u}{|n \times u|}\end{aligned}$$

Viewing

- Look-at
 - Once u , v and n are *normalized*
 - The following will position our camera

$$V = RT = \begin{bmatrix} u_x & u_y & u_z & -eye_x u_x - eye_y u_y - eye_z u_z \\ v_x & v_y & v_z & -eye_x v_x - eye_y v_y - eye_z v_z \\ n_x & n_y & n_z & -eye_x n_x - eye_y n_y - eye_z n_z \\ 0 & 0 & 0 & 1 \end{bmatrix}$$


Viewing

- Look-at
 - Works reasonably well for positioning.
 - But not so well for moving the camera smoothly.

$$V = RT = \begin{bmatrix} u_x & u_y & u_z & -eye_x u_x - eye_y u_y - eye_z u_z \\ v_x & v_y & v_z & -eye_x v_x - eye_y v_y - eye_z v_z \\ n_x & n_y & n_z & -eye_x n_x - eye_y n_y - eye_z n_z \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Viewing

- Yaw, pitch and roll (like an airplane)
 - Here we, essentially, use polar coordinates.
 - A simplified version uses just *azimuth* and *elevation*.
 - Rotate camera in the direction we desire.
 - Translate camera to *eye* point.




Viewing

- Yaw, pitch and roll
 - In reality we perform the *inverse* of what we want.
 - We are transforming world coordinates (all objects) into camera coordinates.
 - Move the world to the camera.
 - » That is, if I want to appear to rotate left 20 degrees.
 - » The transformation about the y -axis would be -20.
 - » Similarly, if I want to appear to move forward 5 units.
 - » I would transform everything by -5.
- So far we have only gotten things in front of the camera!

Viewing

- Projections – Parallel (orthographic)
 - Once in camera coordinates we need a projection transformation to get us to clip coordinates.
 - The transformation matrix that gives us an orthographic projection is:

$$M = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$




Viewing

- Projections – Parallel (orthographic)
 - However, M is applied in the hardware *after* the vertex shader.
 - Which is in clip coordinates
 - How do we “include” or “see” more of our scene?

Viewing

- Projections – Parallel (orthographic)
 - We *scale* what we want to “include” to fit within the canonical view volume. i.e. $(-1,1),(-1,1),(-1,1)$
 - OpenGL provides a function for this called
 - `glOrtho(left, right, bottom, top, near, far)`



Viewing





- Projections – Parallel (orthographic)
 - If you think about it `Ortho` contains a scale and translation.
 - Here is what the transformation matrix looks like.

$$N = ST = \begin{bmatrix} \frac{2}{right - left} & 0 & 0 & -\frac{left + right}{right - left} \\ 0 & \frac{2}{top - bottom} & 0 & -\frac{top + bottom}{top - bottom} \\ 0 & 0 & -\frac{2}{far - near} & -\frac{far + near}{far - near} \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Viewing

- Projections – Perspective
 - Basic symmetrical perspective projection
 - The point (x, y, z) is projected through the projection plane to the eye point (or center of projection COP)
 - We can compute the point of intersection with

$$x_p = \frac{x}{z/d}, \quad y_p = \frac{y}{z/d}$$



Viewing

- Projections – Perspective
 - The simple perspective projection matrix is

$$M = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 1/d & 0 \end{bmatrix}$$

- The important thing to notice here is the position of the $1/d$ term.
 - This means our homogeneous coordinate, w , can be modified (will no longer be 1) when a vertex is multiplied by M .

Viewing

- Projections – Perspective
 - Uh oh, the homogeneous coordinate is no longer 1?


$$q = \begin{bmatrix} x \\ y \\ z \\ z/d \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 1/d & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

- Not the end of the world, remember
 - We have to divide by the homogeneous coordinate to get back to 3D space.

$$q' = \begin{bmatrix} x_p \\ y_p \\ z_p \\ 1 \end{bmatrix} = \begin{bmatrix} \frac{x}{z/d} \\ \frac{y}{z/d} \\ d \\ 1 \end{bmatrix}$$


Viewing

- Projections – Perspective
 - View volume is a pyramid with apex its at the COP.
 - Top and bottom are the near and far clip planes, respectively.
 - Notice that the view and clip planes do not necessarily need to be the same/coincident.





Viewing

- Projections – Perspective
 - OpenGL provides a function, similar to `Ortho`, called `Frustum` with the same parameters.
 - The edges specify the near clip plane.
 - The edges implicitly define the *angle of view* of the projection.



Viewing

- Projections – perspective
 - OpenGL provides a utility function
 - `Perspective(fovy, aspect, near, far)`
 - This form is sometimes more convenient to specify.
 - *Aspect* is the aspect ratio of the view volume.
 - *i.e. width / height*



Viewing

- Projections – perspective
 - Once again, what we had is a projection not the full transformation we need into clip coordinates.
 - The full matrix we do need is defined by:

$$P = NSH = \begin{bmatrix} \frac{right}{near} & 0 & 0 & 0 \\ 0 & \frac{near}{top} & 0 & 0 \\ 0 & 0 & \frac{-far + near}{far - near} & \frac{-2 far \cdot near}{far - near} \\ 0 & 0 & -1 & 0 \end{bmatrix}$$

Viewing

- Projections – perspective
 - Matrices are passed to the vertex shader just like we have seen earlier.
 - Matrices are **uniform** variables where all the data parallel processors on the GPU will see the same value.

```
in vec4 vPosition;  
uniform mat4 modelView;  
uniform mat4 projection;  
  
void main( )  
{  
    //  
    // The perspective division actually happens to gl_Position immediately  
    // after the vertex shader completes. i.e. divided by gl_Position.w  
    //  
    gl_Position = projection * modelView * vPosition;  
}
```